

Assignment No 4

```
import pandas as pd
import numpy as np
import io
import seaborn as sns
import matplotlib.pyplot as plt

%matplotlib inline

from google.colab import files
uploaded=files.upload()

<IPython.core.display.HTML object>

Saving housing.csv to housing.csv

df=pd.read_csv(io.BytesIO(uploaded['housing.csv']))

print(df)
```

	RM	LSTAT	PTRATIO	MEDV
0	6.575	4.98	15.3	504000
1	6.421	9.14	17.8	453600
2	7.185	4.03	17.8	728700
3	6.998	2.94	18.7	701400
4	7.147	5.33	18.7	760200
..
484	6.593	9.67	21.0	470400
485	6.120	9.08	21.0	432600
486	6.976	5.64	21.0	501900
487	6.794	6.48	21.0	462000
488	6.030	7.88	21.0	249900

```
[489 rows x 4 columns]

df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 489 entries, 0 to 488
Data columns (total 4 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   RM          489 non-null    float64
 1   LSTAT       489 non-null    float64
 2   PTRATIO     489 non-null    float64
 3   MEDV        489 non-null    int64
dtypes: float64(3), int64(1)
memory usage: 15.4 KB
```

```
print(df.isnull().sum())
```

```
RM          0
LSTAT       0
PTRATIO     0
MEDV        0
dtype: int64
```

```
print(df.isnull().sum().sum())
```

```
0
```

```
print(np.shape(df))
```

```
(489, 4)
```

```
print(df.describe())
```

	RM	LSTAT	PTRATIO	MEDV
count	489.000000	489.000000	489.000000	4.890000e+02
mean	6.240288	12.939632	18.516564	4.543429e+05
std	0.643650	7.081990	2.111268	1.653403e+05
min	3.561000	1.980000	12.600000	1.050000e+05
25%	5.880000	7.370000	17.400000	3.507000e+05
50%	6.185000	11.690000	19.100000	4.389000e+05
75%	6.575000	17.120000	20.200000	5.187000e+05
max	8.398000	37.970000	22.000000	1.024800e+06


```
#set the size of the figure
```

```
sns.set(rc={'figure.figsize':(11.7,8.27)})
```

```
#Histogram for distribution of the target values
```

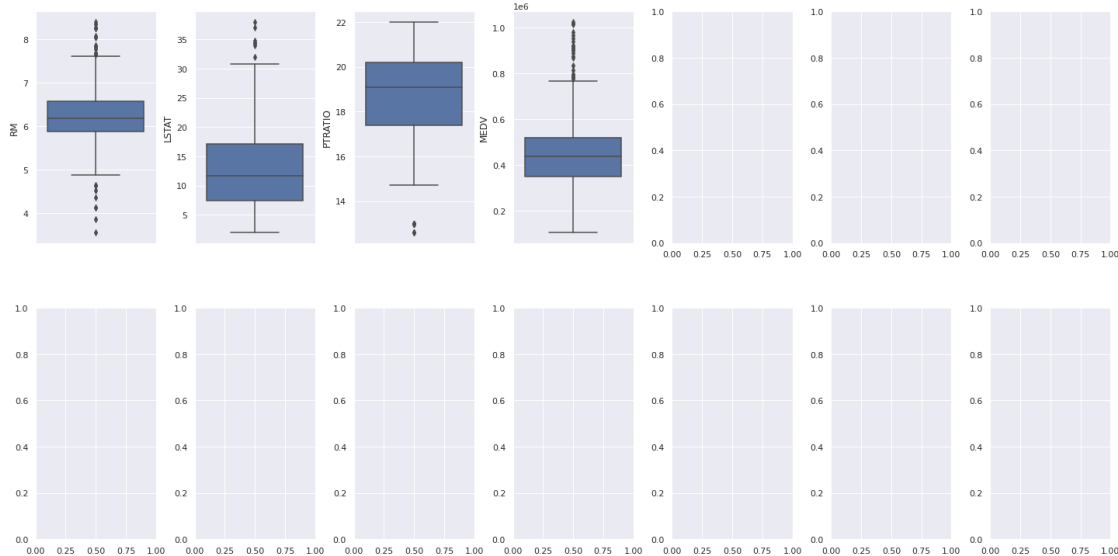
```
sns.distplot(df['MEDV'], bins =30)
plt.show()
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:2619:
FutureWarning: `distplot` is a deprecated function and will be removed in a
future version. Please adapt your code to use either `displot` (a figure-
level function with similar flexibility) or `histplot` (an axes-level
function for histograms).
  warnings.warn(msg, FutureWarning)
```



```
import seaborn as sns
import matplotlib.pyplot as plt
from scipy import stats

fig, axs = plt.subplots(ncols=7, nrows=2, figsize=(20, 10))
index = 0
axs = axs.flatten()
for k,v in df.items():
    sns.boxplot(y=k, data=df, ax=axs[index])
    index += 1
plt.tight_layout(pad=0.4, w_pad=0.5, h_pad=5.0)
```



```
fig, axs = plt.subplots(ncols=7, nrows=2, figsize=(20, 10))
index = 0
axs = axs.flatten()
for k,v in df.items():
    sns.distplot(v, ax=axs[index])
    index += 1
plt.tight_layout(pad=0.4, w_pad=0.5, h_pad=5.0)
```

/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:2619:
FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

warnings.warn(msg, FutureWarning)

/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:2619:
FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).


warnings.warn(msg, FutureWarning)

/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:2619:
FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

warnings.warn(msg, FutureWarning)


/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:2619:
FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

warnings.warn(msg, FutureWarning)



```
correlation_matrix = df.corr().round(2)
sns.heatmap(data=correlation_matrix, annot = True)
```

<matplotlib.axes._subplots.AxesSubplot at 0x7f9d65c7b450>





```
plt.figure(figsize=(20,5))
```

```
features = ['LSTAT', 'RM']
```

```
target = df['MEDV']
```

```
for i, col in enumerate(features):  
    plt.subplot(1, len(features), i+1)  
    x=df[col]  
    y = target  
    plt.scatter(x, y,marker='o')  
    plt.title(col)  
    plt.xlabel(col)  
    plt.ylabel('MEDV')
```



```
#Prepare data for training
```

```
X = pd.DataFrame(np.c_[df['LSTAT'],df['RM']], columns = ['LSTAT','RM'])  
Y = df['MEDV']
```

```
#split data into training and testing sets
```

```
from sklearn.model_selection import train_test_split  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size =  
0.3, random_state = 42)  
print(X_train.shape)  
print(X_test.shape)  
print(y_train.shape)  
print(y_train.shape)
```

```
(342, 2)
```

```
(147, 2)
```

```
(342,)
```

```
(342,)
```

```
from sklearn import linear_model
```

```
# Train the model using sklearn linear regression
```

```
from sklearn.linear_model import LinearRegression  
from sklearn.metrics import mean_squared_error, r2_score
```

```
lin_model = LinearRegression()
```

```
lin_model.fit(X_train,y_train)
```

```
LinearRegression()
```

```
# Model evaluation for training set
```

```
y_train_predict = lin_model.predict(X_train)
```

```
rmse = (np.sqrt(mean_squared_error(y_train, y_train_predict)))  
r2 = r2_score(y_train,y_train_predict)
```

```
print("Model performance for training set")  
print('RMSE is {}'.format(rmse))  
print('R2 score is {}'.format (r2))  
print("\n")
```

```
# Model evaluation for testing set
```

```
y_test_predict = lin_model.predict(X_test)  
rmse = (np.sqrt(mean_squared_error(y_test, y_test_predict)))  
r2 = r2_score(y_test,y_test_predict)
```

```
print("Model performance for testing set")  
print('RMSE is {}'.format(rmse))  
print('R2 score is {}'.format (r2))  
print("\n")
```

```
Model performance for training set  
RMSE is 97613.15525868809  
R2 score is 0.6701982599508145
```

```
Model performance for testing set  
RMSE is 92082.48814292479  
R2 score is 0.628271585004582
```

```
# Plotting y_test vs y_pred
```

```
plt.scatter(y_test, y_test_predict)  
plt.show()
```

